

Topics in Probabilistic Modeling and Inference
(CS698X, Spring 2021)
Homework 3
Due Date: May 1, 2021 (11:59pm)

Instructions:

- Only electronic submissions will be accepted. Your main PDF writeup must be typeset in LaTeX (please also refer to the “Additional Instructions” below).
- Your submission will have two parts: The main PDF writeup (to be submitted via Gradescope <https://www.gradescope.com/>) and the code for the programming part (to be submitted via this Dropbox link: <https://tinyurl.com/67ddj365>). Both parts must be submitted by the deadline to receive full credit (**delay in submitting either part would incur late penalty for both parts**). We will be accepting late submissions upto 72 hours after the deadline (with every 24 hours delay incurring a 10% late penalty, applied on per-hour delay basis). We won't be able to accept submissions after that.
- We have created your Gradescope account (you should have received the notification). Please use your IITK CC ID (not any other email ID) to login. Use the “Forgot Password” option to set your password.

Additional Instructions

- We have provided a LaTeX template file `hw3sol.tex` to help typeset your PDF writeup. There is also a style file `pmi.sty` that contain shortcuts to many of the useful LaTeX commands for doing things such as boldfaced/calligraphic fonts for letters, various mathematical/greek symbols, etc., and others. Use of these shortcuts is recommended (but not necessary).
- Your answer to every question should begin on a new page. The provided template is designed to do this automatically. However, if it fails to do so, use the `\clearpage` option in LaTeX before starting the answer to a new question, to *enforce* this.
- While submitting your assignment on the Gradescope website, you will have to specify on which page(s) is question 1 answered, on which page(s) is question 2 answered etc. To do this properly, first ensure that the answer to each question starts on a different page.
- Be careful to flush all your floats (figures, tables) corresponding to question n before starting the answer to question $n + 1$ otherwise, while grading, we might miss your important parts of your answers.
- Your solutions must appear in proper order in the PDF file i.e. solution to question n must be complete in the PDF file (including all plots, tables, proofs etc) before you present a solution to question $n + 1$.
- For the programming part, all the code and README should be zipped together and submitted as a single file named `yourrollnumber.zip`. Please DO NOT submit the data provided.

Problem 1: This is Bayes Rule! (15 marks)

Given a bunch of observations $\mathbf{x}_1, \dots, \mathbf{x}_N$ drawn i.i.d. from an observation/likelihood model $p(\mathbf{x}|\theta)$, and assuming a prior distribution $p(\theta)$ on the model parameters θ , prove that solving the following problem is equivalent to the Bayes rule for finding the posterior distribution of θ :

$$\arg \min_{q(\theta)} - \sum_{n=1}^N \left[\int q(\theta) \log p(\mathbf{x}_n|\theta) d\theta \right] + KL(q(\theta)||p(\theta))$$

Also give an brief, intuitive explanation of the above objective function (maximum 1-2 sentences).

Problem 2: Mean-Field VI for Sparse Bayesian Linear Regression (30 marks)

Assume N observations $\{\mathbf{x}_n, y_n\}_{n=1}^N$ generated from a regression model $y_n \sim \mathcal{N}(y_n|\mathbf{w}^\top \mathbf{x}_n, \beta^{-1})$. Further assume a Gaussian prior on \mathbf{w} with different component-wise precisions, i.e., $p(\mathbf{w}) = \mathcal{N}(\mathbf{w}|0, \text{diag}(\alpha_1^{-1}, \dots, \alpha_D^{-1}))$. Also assume gamma priors on the noise precision β and prior's precisions $\{\alpha_d\}_{d=1}^D$, i.e., $\beta \sim \text{Gamma}(\beta|a_0, b_0)$ and $\alpha_d \sim \text{Gamma}(\alpha_d|e_0, f_0), \forall d$. We will use the following parametrization of the gamma: $\text{Gamma}(\eta|\tau_1, \tau_2) = \frac{\tau_2^{\tau_1}}{\Gamma(\tau_1)} \eta^{\tau_1-1} \exp(-\tau_2 \eta)$.

Derive the mean-field variational inference algorithm for approximating the posterior distribution

$$q(\mathbf{w}, \beta, \alpha_1, \dots, \alpha_D) = q(\mathbf{w})q(\beta)q(\alpha_1) \dots q(\alpha_D) \approx p(\mathbf{w}, \beta, \alpha_1, \dots, \alpha_D|\mathbf{y}, \mathbf{X})$$

You may use the “recognition” method for inferring each q distribution, or explicitly write down the ELBO for this model and take derivatives w.r.t. the variational parameters of each q distribution to estimate these.

Problem 3: Gibbs Sampling (20 marks)

Suppose we are given a bunch of count-valued observations x_1, \dots, x_N , assumed generated from the following hierarchical model: $p(x_n|\lambda_n) = \text{Poisson}(x_n|\lambda_n)$, $p(\lambda_n|\alpha, \beta) = \text{Gamma}(\lambda_n|\alpha, \beta)$, $n = 1, \dots, N$, $p(\alpha|a, b) = \text{Gamma}(\alpha|a, b)$, and $p(\beta|c, d) = \text{Gamma}(\beta|c, d)$. Assume a, b, c, d to be fixed.

We would like to do Gibbs sampling for this model. To do so, derive the conditional posterior (CP) of each variable $\lambda_1, \dots, \lambda_N, \alpha$, and β , given all the other variables. Are all CPs available in closed form?

Problem 4: Using Samples for Prediction (10 marks)

Consider a matrix factorization model for a partially observed $N \times M$ matrix \mathbf{R} , where $p(r_{ij}|\mathbf{u}_i, \mathbf{v}_j) = \mathcal{N}(r_{ij}|\mathbf{u}_i^\top \mathbf{v}_j, \beta^{-1})$, and \mathbf{u}_i and \mathbf{v}_j denote the latent factors of i -th row and j -th column of \mathbf{R} , respectively. The posterior predictive distribution of each r_{ij} is defined as $p(r_{ij}|\mathbf{R}) = \int p(r_{ij}|\mathbf{u}_i, \mathbf{v}_j)p(\mathbf{u}_i, \mathbf{v}_j|\mathbf{R})d\mathbf{u}_i d\mathbf{v}_j$, which is in general intractable. Suppose we are given a set of S samples $\{\mathbf{U}^{(s)}, \mathbf{V}^{(s)}\}_{s=1}^S$ generated by a Gibbs sampler for this matrix factorization model, where $\mathbf{U}^{(s)} = \{\mathbf{u}_i^{(s)}\}_{i=1}^N$ and $\mathbf{V}^{(s)} = \{\mathbf{v}_j^{(s)}\}_{j=1}^M$.

Given these samples, derive the expressions for the sample based approximation of the *mean (expectation)* as well as the *variance* of any entry r_{ij} of the matrix \mathbf{R} .

Hint: Note that we can write each r_{ij} as $\mathbf{u}_i^\top \mathbf{v}_j + \epsilon_{ij}$ where $\epsilon_{ij} \sim \mathcal{N}(\epsilon_{ij}|0, \beta^{-1})$.

Problem 5: Rejection Sampling (25 marks)

Consider a distribution $p(x) \propto \exp(\sin(x))$ for $-\pi \leq x \leq \pi$. Denote $\exp(\sin(x))$ as $\tilde{p}(x)$. Suppose we want to use Rejection Sampling to sample from $p(x)$ and use a proposal distribution $q(x) = \mathcal{N}(x|0, \sigma^2)$. Find out the expression for the optimal value of the constant M such that $Mq(z) \geq \tilde{p}(x)$, as required in Rejection Sampling.

Using this value of M and some suitably chosen σ^2 , draw 10,000 samples from $p(x)$ and plot the resulting histogram of the samples. Submit your code in form of a Python notebook.